



Programming the Continuum

*Towards Better Techniques for Developing
Distributed Science Applications*

Dissertation Defense
by J. Gregory Pauloski


Committee: Kyle Chard (advisor), Ian Foster (advisor), Michael Franklin

Why Program the Continuum?

Better and More Ambitious Science

Computing continuum: cyberinfrastructure spanning edge devices, the cloud, and supercomputers

- Faster & more reliable networks
- Specialized accelerators
- Data locality
- Performance requirements
- Compute availability & costs
- Better cloud management




Imagine you are a computational scientist...


With a distributed science application to build...

What framework do you use?

What resources will I use?



Provisioned or Serverless?



Strong Ecosystem

Weak Ecosystem



Challenges in Programming the Continuum

Distribute computational tasks across federated devices? Possible.

- Globus Compute distributed FaaS Model

Manage intermediate data between tasks? Limited.

- Interoperability between distributed/parallel frameworks is challenging
- Cloud object storage is reliable/available but expensive for data-intensive apps
- P2P CDNs are good for edge devices but bad for clusters

Build persistent and loosely coupled components? Limited.

- Easy in cloud-native apps (microservice architectures)
- Hard in federated apps (requires ad-hoc solutions)

Programming the Continuum

New programming techniques **enable and accelerate** task-centric **science applications** executed **across the computing continuum**.

P1	What are the limitations in existing distributed computing frameworks?	eScience '24 (Best Paper)
P2	How to represent and efficiently move objects across federated systems?	SC '23 & HPPSS '24
P3	How to support common high-level data flow patterns?	TPDS '24
P4	How to build and deploy stateful agents across federated systems?	IEEE Computer* & SC '25*

Better, easier, & faster science! — MLHPC '21, IJHPCA '23, HCW '23, IJHPCA '24, CCGRID '25
& Others In Review/Progress

*Under Review / In Progress



Task Performance Suite



ProxyStore



Proxy Patterns



Federated Agents

Modern Science Applications are *Task-centric*

Applications are composed as a set of **discrete tasks** designed to **automate** computational processes to achieve a **scientific goal**

Benefits

- Heterogeneous Resources
- Software Modularity
- Monitoring
- Performance
- Reproducibility
- *and many more!*

Applications [1]

- Bioinformatics
- Cosmology
- High Energy Physics
- Materials Science
- Molecular Dynamics
- *and many more!*

Challenges [2]

- Coupling AI/ML/Quantum
- Cloud and HPC Integration
- Data Flow/Provenance
- Standards/Interoperability
- Performance
- *and many more!*

[1] "Scientific Workflows: Moving Across Paradigms" (<https://dl.acm.org/doi/10.1145/3012429>)

[2] Workflows Community Summit (<https://arxiv.org/abs/2304.00019>)

Task Execution Frameworks

Manage the execution of tasks in parallel across arbitrary hardware.

Workflow Management Systems

Define, manage, and execute workflows represented by a directed acyclic graph (DAG) of tasks

Explicit

DAG defined via configuration file or domain specific language



Implicit

Task dependencies derived through dynamic evaluation of a procedural script



Concurrent Executors

On-demand asynchronous execution of tasks




How do we **benchmark** and compare execution frameworks?

TaPS: Task Performance Suite

- Reference set of applications to standardize benchmarking workloads
- Robust and reproducible configuration system
- Benchmark task executors & data management systems

Guide future research!




<https://taps.proxystore.dev/latest/api/>

Applications: *Benchmarking Workloads*

- Seven Real Apps
- Two Synthetic
- Diverse Patterns
- Diverse Domains
- Per-App Guides
- Add your own!

Type	Name	Domain	Task Type(s)	Data Type(s)
Real	cholesky	Linear Algebra	Python	In-memory
	docking	Drug Discovery	Executable, Python	File
	fedlearn	Machine Learning	Python	In-memory
	mapreduce	Text Analysis	Python	File, In-memory
	moldesign	Molecular Design	Python	In-memory
	montage	Astronomy	Executable	File
	physics	Mechanics	Python	In-memory
Synthetic	synthetic	—	Python	In-memory
	failures	—	Depends on base app	Depends on base app




<https://taps.proxystore.dev/latest/apps/>

What did we learn?

Many things!

But most *important* to this story is...

- No single executor is the best at everything.
- Large-scale federated apps will need to use multiple concurrently for optimal results.



*Task data exceeds Globus Compute 10 MB payload limit.

<https://github.com/proxystore/escience24-taps-analysis>

Task Performance Suite

ProxyStore

Proxy Patterns

Federated Agents

Representing Intermediate Objects

In a federated environment, how do we...

- **Represent** an object x such that the producer and consumers of x can globally reference x ?
 - ◆ Assume x is immutable in the context of intermediate objects
- **Communicate** x from producer to consumers when consumers
 - ◆ are not known ahead of time,
 - ◆ can be located in different places, and
 - ◆ have different optimal communication methods?


Representing Intermediate Objects

Case Study: Ray

- Ray **represents** x with an object ref
 - ✓ Distributed reference counting
 - ✓ Cheap to pass around
 - ✗ Not valid outside of the Ray cluster it was created in
- Ray **communicates** x using RPCs
 - ✓ Fast & direct within a cluster
 - ✗ RPC not possible outside of cluster



ProxyStore



Data flow management library for distributed Python workflows

- Represent and efficiently move objects in federated applications
- Proxy **transparently** decouples control and data flow
- Best of both **pass-by-reference** and **pass-by-value**
- Use any mediated communication method via plugins

Proxy Objects

What is a proxy (in this context)?

- Self-contained wide-area **reference** to a **target** object
- Transparently resolve target **just-in-time** when first used

What are the benefits?


- Performance (pass-by-reference, async resolve, skip unused objects)
- Reduce code complexity
- Partial resolution of complex objects
- Access control

```
from proxystore.connectors import RedisConnector
from proxystore.store import Store
from proxystore.proxy import Proxy

def foo(x: Bar) -> ...:
    # Resolve of x deferred until use
    assert isinstance(x, Bar)
    # More computation...

with Store('demo', RedisConnector(...)) as store:
    x = Bar(...)
    p = store.proxy(x) # Anything can be proxied
    assert isinstance(p, Proxy)
    foo(p) # Proxies can be passed-by-ref anywhere
```

(2) Store gives object to Connector and generates a Proxy with metadata/Factory.




Connectors

- Comprehensive **mediated** methods (producer/consumer may be temporally decoupled)
- Connector = Python **Protocol**
- **MultiConnector**: Policy-based routing between instances


Protocol	Storage	Intra-Site	Inter-Site	Persistence
File System	Disk	✓		✓
Redis/KeyDB	Hybrid	✓		✓
Margo	Memory	✓		
UCX	Memory	✓		
ZMQ	Memory	✓		
Globus	Disk		✓	✓
DAOS	Disk*	✓		✓
P2P Endpoint	Hybrid	✓	✓	✓

Where do we use ProxyStore?


Multi-Site Workflows




Federated Learning




Real-time Processing




Configuration	Proxied	Time (ms)	Improvement
Globus Compute baseline	—	3411 ± 389	—
FileStore	Inputs	2318 ± 130	32.1%
	Inputs/Outputs	2160 ± 46	36.6%
EndpointStore	Inputs	2375 ± 98	30.4%
	Inputs/Outputs	2280 ± 107	33.2%



Reduce Scheduler Overhead





RDMA in FaaS Systems




P2P Networking

P2P Endpoints: Easy* Multi-Site Workflows

Moving data between sites through the cloud is impractical!



ProxyStore Endpoints: Move proxies through the cloud and data peer-to-peer with UDP hole punching




```
$ proxystore-endpoint configure demo --relay wss://relay.proxystore.dev  
$ proxystore-endpoint start demo # Runs as a daemon process
```

* Easy = no SSH tunnels, one-time setup, no cloud fees

docs.proxystore.dev/main/guides/endpoints/

Multi-site Active Learning


Science Goal: Use quantum chemistry simulations and surrogate ML models to efficiently identify electrolytes with high ionization potentials in a candidate set.



Logan Ward, J. Gregory Pauloski, Valerie Hayot-Sasson, Ryan Chard, Yadu Babuji, Ganesh Sivaraman, Sutanay Choudhury, Kyle Chard, Rajeev Thakur, and Ian Foster. Cloud services enable efficient AI-guided simulation workflows across heterogeneous resources. In Heterogeneity in Computing Workshop at IPDPS. IEEE Computer Society, 2023.

Multi-site Active Learning

Systems Goal: Reduce task communication overheads in workflow system to increase system utilization and task throughput.



Baseline: Parsl manages all intermediate data (transfer via manually created SSH channels)

ProxyStore: MultiConnector

- *Simulation:* Redis
- *Training:* P2P Endpoints
- *Inference:* Globus Transfer/ P2P Endpoints

Takeaways

- Reduce overheads in task scheduler
- Reduce communication of re-used data
- Optimize communication method per data type
- No changes to task code needed

Task Performance Suite

ProxyStore

Proxy Patterns

Federated Agents


Yet...

Object proxy is a *low-level* paradigm:

- A great building block within larger frameworks
- Has known limitations

What are *higher-level* proxy patterns?

- Accelerate development of more sophisticated applications
- Address limitations




1. ProxyFutures

Futures in {Dask, Parsl, Ray, ...}

- Control & data synchronization tightly coupled (no optimization)
- Transfer mechanism fixed
- Not usable outside framework


ProxyFutures

- Explicit & Implicit Usage
- Data synchronization only (good)
- Any transfer mechanism
- Framework-agnostic



1. ProxyFutures

36% reduction
in makespan!




1000 Genomes executed using **Globus Compute** (no task data dependency support) on Chameleon Cloud

2. ProxyStream

High-performance stream processing


- Common in scientific computing
- Data are very large (suboptimal for Kafka-like systems)
- Quickly (1) decide if data should be used and (2) dispatch to node in cluster (e.g., for simulation)



- ProxyStream decouples metadata from bulk data transfer
- Send proxies + metadata through message broker
- Resolve proxies only when needed via more performant methods

2. ProxyStream

Performance equal to or better than state-of-the-art



Synthetic scaling test

- Stream process & dispatch
- Random data & simulated compute

DeepDriveMD

32% faster inference times
21% more inferences done



Use ProxyStream for stateful ML inference workers in pure-functional frameworks

3. Proxy Ownership

Memory Management with Proxies

- No reference counting—distributed reference counting in a federated environment is challenging
- Freeing a proxy can cause errors in other processes sharing the proxy (essentially a null-pointer exception)
- Forgetting to free can cause memory leaks

ProxyStore provides guidance on handling these but it's ultimately up to user

3. Proxy Ownership

Map scope of proxies to tasks.

- Child tasks can borrow a proxy from parent
- Borrowed proxy is valid for tasks' lifetime
- Out-of-scope proxies deleted
- `StoreExecutor` for easy integration of ownership with execution frameworks

Custom lifetimes for more complex scenarios.

- Code-segment, time-leased, and static lifetimes
- Extensible—create your own lifetime types

Rust Ownership/Borrowing Inspired 🦀

Enforced at Runtime

Ownership Rules

1. Each object in the `Store` has an associated `OwnedProxy`
2. There can only be one `OwnedProxy` for any object
3. When the `OwnedProxy` goes out of scope the object is deleted

Bonus details!

Reference Rules

1. At any given time, an `OwnedProxy` may be mutable borrowed once or immutably borrowed many times
2. An `OwnedProxy` and `RefMutProxy` are mutable references; a `RefProxy` is an immutable reference




Last 2.5 years

Since candidacy (~4 months)

Autonomous discovery “harnesses the power of robotics, ML, and AI to **solve big problems** [...] **faster than ever before.**”

Credit: ANL, “[Science 101: Autonomous Discovery](#)”

Challenge 1: Complexity is a Barrier



Challenge 2: Humans are a Bottleneck

Humans synthesize knowledge and propose hypotheses

Humans write, debug, and run programs

Humans interpret results to inform new hypotheses

Agents can be the driving entities

- Persistent, stateful, cooperative
- Intermittent human oversight

Inefficient use of research infrastructure

We need to be here

Credit: Ian Foster, "Empowering Science with Intelligent Middleware and Embodied Agents"

Solution: Multi-Agent Systems for Science

- ✓ Automate closed-loop processes
- ✓ Natural expression of scientific resources (compute, instruments, repositories)
- ✓ Operate autonomously but still cooperatively
- ✓ Execute multi-stage computational science processes
- ✓ Reduce mundane task responsibilities of scientists

The whole is greater than the sum of its parts.
- Aristotle

How do we build agents?

We are missing the middleware to build and connect our agents!


Search database, invoke code, query LLM, ...

A **computational system** that can **interact** with its **environment** and **learn** from those interactions

Data repositories, HPC, robotic labs, other agents


Accumulate data, adapt processes, improve answers

Credit: Ian Foster, "Empowering Science with Intelligent Middleware and Embodied Agents"



Middleware Open Challenges

- Access & privileges
 - Agent discovery
 - Asynchronous communication
 - Fault tolerance
 - Interfaces
 - Mobility
 - Persistent stateful execution
 - Provenance
 - Many more...
- Areas we focused on...*



Under review in IEEE Computer

What does the middleware look like?

Workflows

Dask, Parsl, Pegasus

++ Task automation

++ Distributed task execution

Actor Systems

Akka, Dask, Ray

++ Stateful computation

++ Actor-to-actor interaction

Function-as-a-Service


Globus Compute, Lambda

++ Remote execution

++ Fire-and-forget model

Academy

- *Fire-and-forget*: Agents spawned across remote/federated resources
- *Autonomy*: Agents have agency over their actions and local state/resources
- *Cooperative*: Agents interact to execute tasks & workflows



Communication & Execution

Exchange

- Asynchronous communication through mailboxes
- Every agent/client in system has a unique mailbox
- Local & distributed implementations
- Optimized for low-latency
- Hybrid communication model
- Prefer direct communication between agents when possible; fall back to indirect communication via object store
- Pass-by-reference with ProxyStore for large data

Launcher

- Not required but enables remote execution of agents
- Returns handle to launched agent
- Local threads or processes
- Distributed with Parsl
- Federated with Globus Compute

Agents defined by a behavior

Clients & other agents can request actions

```
import time, threading
from academy.behavior import Behavior, action, loop

class Example(Behavior):
    def __init__(self) -> None:
        self.count = 0 # State stored as attributes

    @action
    def square(self, value: float) -> float:
        return value**2

    @loop
    def count(self, shutdown: threading.Event) -> None:
        while not shutdown.is_set():
            self.count += 1
            time.sleep(1)
```

Instance of a behavior is state

Control loops for autonomous behavior

Single interface
for managing
your agents

Interact with
agents via
handles

```
from academy.exchange.thread import ThreadExchange
from academy.launcher.thread import ThreadLauncher
from academy.manager import Manager

with Manager(
    exchange=ThreadExchange(), # Can be swapped
    launcher=ThreadLauncher(),
) as manager:
    behavior = Example() # From the prior slide
    handle = manager.launch(behavior)
    future = handle.square(2)
    assert future.result() == 4

    handle.shutdown() # Or via the manager
    manager.shutdown(handle.agent_id, blocking=True)
```

Choose exchange
& launcher for
environment

Pass handles to
other agents

Features (rapid fire)



- Any number of actions & control loops
- Special purpose control loop decorators
- Multi-threaded/non-blocking action execution
- Startup and shutdown callbacks
- State persistence plugins
- Re-execution on failure
- Agents can launch other agents
- Discovery/lookup based on behavior
- Handle mailbox multiplexing



*Any interesting? Ask
about them at the end!*


Comparisons to Actor Systems

Why we need ProxyStore!



Low-memory overhead

Low-latency messaging




Experiments performed on Aurora @ ALCF

Use Case: MOF Discovery

Metal Organic Frameworks (MOF)


- Composed of organic molecules (ligands) and inorganic metals (nodes)
- The sponges of materials science!
- Porous structures that adsorb and store gases
- Topologies can be optimized for targeted gas storage → **Carbon Capture**

How to efficiently discover MOFs with desirable properties for target applications?




Intractable search space of ligand, node, & geometry combinations

MOFA: Online learning + GenAI + Simulation




Yan et al., “**MOFA: Discovering Materials for Carbon Capture with a GenAI- and Simulation-Based Workflow**” (Under Review)

MOFA through Autonomous Agents



MOFA Agents Trace





Why is this agentic model better?

- **Placement:** Move agents to resources
- **Separation of concerns:** Resource acquisition and scaling based on local workload
- **Loose coupling:** Swap agents or integrate new agents (e.g., SDL)
- **Shared agents:** Multiple workflows can share agents (microservice-like)

Summary

Impact



Empowering large-scale science through open-source software

Summary

New programming techniques **enable and accelerate** task-centric **science applications** executed **across the computing continuum**.

P1	TaPS : Support research in distributed/parallel execution	eScience '24 (Best Paper)
P2	ProxyStore : Better object references for federated environments	SC '23 & HPPSS '24
P3	Proxy Patterns : Better data flow patterns with object proxies	TPDS '24
P4	Federated Agents : Build science agents for autonomous discovery	IEEE Computer* & SC '25*

Better, easier, & faster science! — MLHPC '21, IJHPCA '23, HCW '23, IJHPCA '24, CCGRID '25
& Others In Review/Progress

*Under Review / In Progress

Programming the Continuum

Towards Better Techniques for Developing
Distributed Science Applications

New programming techniques enable and
accelerate task-centric science applications
executed across the computing continuum.

- **TaPS** [eScience '24]
- **ProxyStore** [SC '23 & HPPSS '24]
- **Proxy Patterns** [TPDS '24]
- **Federated Agents** [IEEE Computer* & SC '25*]

Better, easier, & faster science!

MLHPC '21, IJHPCA '23, HCW '23, IJHPCA '24,
CCGRID '25 & Others In Review/Progress

Questions?

Contact:

J. Gregory Pauloski

jgpauloski@uchicago.edu

Reference:

github.com/proxystore

docs.proxystore.dev

taps.proxystore.dev



github.com/proxystore

Acknowledgements:

- Argonne National Laboratory under U.S. Department of Energy Contract DE-AC02-06CH11357
- National Science Foundation under Grant 2004894 and Grant 2209919
- ExaLearn Co-design Center of the Exascale Computing Project (17-SC-20-SC)